

**Advanced Topics in Numerical Analysis:  
High Performance Computing  
Assignment 4 (due Apr. 18, 2022)  
Anav Prasad, ap7152**

**Handing in your homework:** Hand in your homework as for the previous homework assignments (git repo with Makefile), answering the questions by adding a text or a  $\text{\LaTeX}$  file to your repo.

1. **Matrix-vector operations on a GPU.** Write CUDA code for an inner product between two given (long) vectors on a GPU. Then, generalize this code to implement a matrix-vector multiplication (no blocking needed here) on the GPU. Check the correctness of your implementation by performing the same computation on the CPU and compare them. Report the memory band your code obtains on different GPUs<sup>1</sup>.

**Solution:**

For **inner-product**, the vector size was chosen as  $2^{24}$  because in the ranges of  $2^1 - 2^{10}$ , the CPU bandwidths and GPU bandwidths are comparable. In the ranges of  $2^{11} - 2^{20}$ , CPU performs better than the GPU. Finally, at sizes  $> 2^{20}$ , GPU parallizations finally become significant enough to improve the overall time as seen below.

Note that the values of bandwidth's are in GB/s.

• **inner-product results:**

Vector Size:  $2^{24}$

| Compute Server | GPU             | CPU Bandwidth | GPU Bandwidth |
|----------------|-----------------|---------------|---------------|
| cuda1          | GTX TITAN Black | 12.0025       | 62.7648       |
| cuda2          | RTX 2080 Ti     | 10.955        | 131.298       |
| cuda3          | TITAN V         | 5.0364        | 204.313       |
| cuda4          | GTX TITAN X     | 5.21459       | 66.8682       |
| cuda5          | GTX TITAN Z     | 8.89145       | 51.611        |

For **matrix-vector-product**, on the other hand, GPU parallelized code almost always performs at least a little better than the CPU parallelized code at varied ranges of matrix and vector sizes. The given matrix size shows decent distinction between the two bandwidths.

---

<sup>1</sup>The cuda{1-5}.cims.nyu.edu compute servers at the Institute have different Nvidia GPUs, for an overview see the list of compute servers available at the Institute: <https://cims.nyu.edu/webapps/content/systems/resources/computeservers>.

- **matrix-vector-product results:**

Matrix Size:  $500 \times 2^{18}$

Vector Size:  $2^{18}$

| Compute Server | GPU             | CPU Bandwidth | GPU Bandwidth |
|----------------|-----------------|---------------|---------------|
| cuda1          | GTX TITAN Black | 13.7252       | 37.0873       |
| cuda2          | RTX 2080 Ti     | 5.11759       | 60.84         |
| cuda3          | TITAN V         | 5.6556        | 90.4507       |
| cuda4          | GTX TITAN X     | 6.3598        | 39.2735       |
| cuda5          | GTX TITAN Z     | 12.341        | 31.9898       |

2. **2D Jacobi method on a GPU.** Implement the 2D Jacobi method as discussed in the 2nd homework assignment using CUDA. Check the correctness of your implementation by performing the same computation on the CPU and compare them.

**Solution:**

Machine Details:

Compute Server cuda3.cims.nyu.edu

Device Name: NVIDIA TITAN V

Computation Capability: 7.0

Memory: 12.6528

Number of Threads Set = 4

| N    | #Max Iterations | CPU Time | GPU Time  |
|------|-----------------|----------|-----------|
| 100  | 10000           | 1.01629s | 0.342811s |
| 1000 | 10000           | 9.07443s | 6.78653s  |

3. **Update on final projection** Describe with a few sentences the status of your final project: What tasks that you formulated in the proposal have you worked on? Did you run into unforeseen issues?

**Solution:**

Collaborators:


- Aditya Kashilkar, ask9126
- Sai Himall Allu, sa6783

The following checkpoints have been completed:


- Checkpoint 1: Benchmark the serialized implementation of WFC algorithm

A few sample images:

– chess.bmp:



**Figure 1:** Reference bitmap:  
chess.bmp




**Figure 2:** Larger candidate  
bitmap generated by the algorithm

Time taken to process and generate the larger bitmap from chess.bmp: 0.504s

– city.bmp:




**Figure 3:** Reference bitmap:  
city.bmp




**Figure 4:** Larger candidate  
bitmap generated by the algorithm

Time taken to process and generate the larger bitmap from city.bmp: 1.928s

– flowers.bmp:



**Figure 5:** Reference bitmap:  
flowers.bmp



**Figure 6:** Larger candidate bitmap generated by the al-  
gorithm

Time taken to process and generate the larger bitmap from flowers.bmp: 1.388s

Total Time taken to process all 54 sample images: 54.671533

4. Checkpoints 2 and 3:

- > Analyze the reference code for optimal parallelization avenues.
- > Implement basic OpenMP optimizations.

Although there are simple optimisation routines possible for the multiple for-loops used in the code-base, there is some amount of non-triviality involved with optimising it.

While the code compiles and, the results are holding up, the speedups reported are not that great. We reason this might be due to the fact that there are a significant number of branching operations in our serial code which is negating some of the benefits of parallelization.

We are currently looking into avenues for reducing the number of branching instructions through some optimisation tricks